

Census Income Project

Introduction :

In this project, we are going to predict whether a person's income is above 50k or below 50k using various features like age, education, and occupation. The dataset we are going to use is the census income dataset which contains about 32560 rows and 15 features .

Defining the problem statement :

The data contains anonymous information such as age, occupation, education, working class, etc. *The goal is to train a binary classifier to predict the income which has two possible values '>50K' and '<50K'.* There are 48842 instances and 14 attributes in the dataset. The data contains a good blend of categorical, numerical and missing values.

The dataset contains the labels which we have to predict which is the dependent feature 'Income level'. This feature is discrete consisting of two categories income less than 50k and more than 50k. So the problem we have is a **Supervised Binary Classification** type.

Data Pre-processing Done:

The purpose of preprocessing is to convert raw data into a form that fits machine learning. Structured and clean data allows a data scientist to get more precise results from an applied machine learning model. The technique includes data formatting, cleaning, and sampling.

- **Hardware and Software Requirements and Tools Used**

- *Tools: Python 3.8.5, Jupyter Notebook, Numpy, Pandas, Matplotlib, Seaborn, Scikit-learn, Scipy*
- *Techniques: LogisticRegression, GaussianNB, SVC, RandomForestClassifier, AdaBoostClassifier, DecisionTreeClassifier, LinearSVC.*
- *Hardware: I3 processor, 4GB RAM*

- **Model/s Development and Evaluation :**

I have used the below models for classification:

1. LogisticRegression
2. SVC
3. RandomForestClassifier
4. AdaBoostClassifier
5. DecisionTreeClassifier
6. *GaussianNB*
7. *LinearSVC*
8. GradientBoostingClassifier
9. MultinomialNB
10. KNeighborsClassifier

- **Identification of possible problem-solving approaches (methods)**

- Read the data (from csv)
- Identify the dependent and independent variables.
- Check if the data has missing values or the data is categorical or not.
- If yes, apply basic data preprocessing operations to bring the data in a go to go format.
- Now split the data into the groups of training and testing for the respective purpose.
- After splitting data, fit it to a most suitable model. (How to find a suitable model is answered below)
- Validate the model. If satisfactory, then go with it, else tune the parameters and keep testing. In a few cases, you can also try different algorithms for the same problem to understand the difference between the accuracies.
- From step 7 one can also learn about accuracy paradox.
- Visualize the data.

- **Testing of Identified Approaches (Algorithms)**

1. Naive Bayes
2. Cross validation
3. Confusion matrix
4. Accuracy score
5. Classification report

- **Run and Evaluate selected models**

I have used the below models for classification:

Logistic regression

```
] : lg=LogisticRegression()  
lg.fit(x_train,y_train)  
pred=lg.predict(x_test)  
lg_accu=accuracy_score(y_test,pred)  
print(accuracy_score(y_test,pred))  
print(confusion_matrix(y_test,pred))  
print(classification_report(y_test,pred))
```

0.7849232201023731

[[7668 443]

[1868 766]]

	precision	recall	f1-score	support
0	0.80	0.95	0.87	8111
1	0.63	0.29	0.40	2634
accuracy			0.78	10745
macro avg	0.72	0.62	0.63	10745
weighted avg	0.76	0.78	0.75	10745

KNeighborsClassifier

```
: from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier()
knn.fit(x_train,y_train)
knn_pred=knn.predict(x_test)
knn_accu=accuracy_score(y_test,pred)

print(accuracy_score(y_test,pred))
print(confusion_matrix(y_test,pred))
print(classification_report(y_test,pred))
```

0.7849232201023731

[[7668 443]

[1868 766]]

	precision	recall	f1-score	support
0	0.80	0.95	0.87	8111
1	0.63	0.29	0.40	2634
accuracy			0.78	10745
macro avg	0.72	0.62	0.63	10745
weighted avg	0.76	0.78	0.75	10745

RandomForestClassifier

```
|: rf=RandomForestClassifier()
rf.fit(x_train,y_train)
pred=rf.predict(x_test)
rf_accu=accuracy_score(y_test,pred)
print(accuracy_score(y_test,pred))
print(confusion_matrix(y_test,pred))
print(classification_report(y_test,pred))
```

0.8538855281526291

[[7520 591]

[979 1655]]

	precision	recall	f1-score	support
0	0.88	0.93	0.91	8111
1	0.74	0.63	0.68	2634
accuracy			0.85	10745
macro avg	0.81	0.78	0.79	10745
weighted avg	0.85	0.85	0.85	10745

SVC

```
: svc=SVC()
  svc.fit(x_train,y_train)
  pred=svc.predict(x_test)
  svc_accu=accuracy_score(y_test,pred)
  print(accuracy_score(y_test,pred))
  print(confusion_matrix(y_test,pred))
  print(classification_report(y_test,pred))
```

0.7879944160074454

[[8101 10]

[2268 366]]

	precision	recall	f1-score	support
0	0.78	1.00	0.88	8111
1	0.97	0.14	0.24	2634
accuracy			0.79	10745
macro avg	0.88	0.57	0.56	10745
weighted avg	0.83	0.79	0.72	10745

AdaBoostClassifier

```
ad=AdaBoostClassifier()
ad.fit(x_train,y_train)
pred=ad.predict(x_test)
ad_accu=accuracy_score(y_test,pred)
print(accuracy_score(y_test,pred))
print(confusion_matrix(y_test,pred))
print(classification_report(y_test,pred))
```

0.8577943229409027

[[7596 515]

[1013 1621]]

	precision	recall	f1-score	support
0	0.88	0.94	0.91	8111
1	0.76	0.62	0.68	2634
accuracy			0.86	10745
macro avg	0.82	0.78	0.79	10745
weighted avg	0.85	0.86	0.85	10745

GaussianNB

```
: gnb=GaussianNB()
  gnb.fit(x_train,y_train)
  pred=gnb.predict(x_test)
  gnb_accu=accuracy_score(y_test,pred)
  print(accuracy_score(y_test,pred))
  print(confusion_matrix(y_test,pred))
  print(classification_report(y_test,pred))
```

0.7892042810609586

[[7665 446]

[1819 815]]

	precision	recall	f1-score	support
0	0.81	0.95	0.87	8111
1	0.65	0.31	0.42	2634
accuracy			0.79	10745
macro avg	0.73	0.63	0.64	10745
weighted avg	0.77	0.79	0.76	10745

MultinomialNB

```
: from sklearn.naive_bayes import MultinomialNB
  mnb=MultinomialNB()
  mnb.fit(x_train,y_train)
  pred=mnb.predict(x_test)
  mnb_accu=accuracy_score(y_test,pred)
  print(accuracy_score(y_test,pred))
  print(confusion_matrix(y_test,pred))
  print(classification_report(y_test,pred))
```

0.7779432294090275

[[7743 368]

[2018 616]]

	precision	recall	f1-score	support
0	0.79	0.95	0.87	8111
1	0.63	0.23	0.34	2634
accuracy			0.78	10745
macro avg	0.71	0.59	0.60	10745
weighted avg	0.75	0.78	0.74	10745

LinearSVC

```
: LSVC=LinearSVC()  
LSVC.fit(x_train,y_train)  
pred=LSVC.predict(x_test)  
LSVC_accu=accuracy_score(y_test,pred)  
print(accuracy_score(y_test,pred))  
print(confusion_matrix(y_test,pred))  
print(classification_report(y_test,pred))
```

0.76379711493718

[[8106 5]

[2533 101]]

	precision	recall	f1-score	support
0	0.76	1.00	0.86	8111
1	0.95	0.04	0.07	2634
accuracy			0.76	10745
macro avg	0.86	0.52	0.47	10745
weighted avg	0.81	0.76	0.67	10745

DecisionTreeClassifier

```
dt=DecisionTreeClassifier()  
dt.fit(x_train,y_train)  
pred=dt.predict(x_test)  
dt_accu=accuracy_score(y_test,pred)  
  
print(accuracy_score(y_test,pred))  
print(confusion_matrix(y_test,pred))  
print(classification_report(y_test,pred))
```

0.8078175895765473

[[7021 1090]

[975 1659]]

	precision	recall	f1-score	support
0	0.88	0.87	0.87	8111
1	0.60	0.63	0.62	2634
accuracy			0.81	10745
macro avg	0.74	0.75	0.74	10745
weighted avg	0.81	0.81	0.81	10745

GradientBoostingClassifier

```
gbc=GradientBoostingClassifier()
gbc.fit(x_train,y_train)
pred=gbc.predict(x_test)
gbc_accu=accuracy_score(y_test,pred)

print(accuracy_score(y_test,pred))
print(confusion_matrix(y_test,pred))
print(classification_report(y_test,pred))
```

0.8586319218241042

[[7664 447]

[1072 1562]]

	precision	recall	f1-score	support
0	0.88	0.94	0.91	8111
1	0.78	0.59	0.67	2634
accuracy			0.86	10745
macro avg	0.83	0.77	0.79	10745
weighted avg	0.85	0.86	0.85	10745

- Conclusion:

conclusion																					
	0	1	2	3	4	5	6	7	8	9	...	10735	10736	10737	10738	10739	10740	10741	10742	10743	10744
predicted	0	0	0	0	0	0	1	0	0	0	...	1	0	0	1	1	0	1	0	0	0
original	0	0	0	0	0	0	1	0	0	0	...	0	0	0	1	1	0	1	0	0	0

2 rows × 10745 columns

GradientBoostingClassifier is accuracy score ,precision ,recall ,CV and f1-score is good than other models,Hence GradientBoostingClassifier is performing is good. GradientBoostingClassifier has accuracy score is 85% . GradientBoostingClassifier is best model for this dataset.